

03:47. The export is already out.

An AI agent calls **crm.bulk_export** with **{email, pesel}**. It holds a password, not an order. Nobody is awake to stop it.

- The log proves **what** happened. It cannot prove **who authorized it**.
- The first signal arrives at 09:15 — the data left hours earlier.
- The regulator asks one question: **under whose authority?**

You cannot trust what you cannot stop.

Trust is not a setting. It is the **ability to revoke**.

Permission today is granted once and lives forever — a key in production. What is missing is not another view. It is **the moment a permission ends**.

- **Four barriers, one cause.** Culturally an agent is not trusted as a colleague; socially it is “if it happens, it is on me”; product-wise there is nothing to stop with; in the story we tell ourselves, “we see everything already”.
- The old way: **a log you read on Monday** — evidence after the fact, never authority before it.
- So the unit is not visibility. It is **authorized action** — and the receipt that proves it.

No approval, no action. **TENET** holds that one rule.

Every tool-call passes the kernel **before it executes**: scope, TTL, signer. One click takes the authority back — and the agent stops mid-flight, not after the export.

AUTHORITY

Identity is bound to one agent and one grant. A grant for one order never authorizes another.

GRANT

HMAC-SHA256 over the scope. The signature is verified **at the gate**, before any rule is read.

REVOKE

One call halts the agent mid-flight. The next call returns REVOKED, not “unknown”.


RECEIPT

Append-only, hash-chained proof: who authorized, under which order, how long it lived.

- Every call is either **rejected before execution** or carries a **receipt**.
- Revocable by design — the order exists to be taken back, not to be filed.

A control layer with **a door** — and it opens only for you.

- 1 you **say** what the agent should do — typed or spoken
- 2 **your AI proposes** one concrete action and says what it will touch
- 3 the **kernel decides**: allowed, not allowed, or waiting for you
- 4 you **decide** — and only your decision moves it; nothing runs on its own
- 5 the action is **witnessed** crossing the line (or **stopped** at it)
- 6 you get the **receipt** — a real record from a real round
- 7 **WIN**: the agent acted **inside the line you drew**



A-0002
the action that crossed the line

HTTP 200 · 19.5 ms
the upstream really answered

receipt 4db4d785
hash-chained, recomputable

read from the live control plane —
not typed into this slide

One **TENET**, in the film and in the product.

- One **locked studio voice** for the whole cycle — eight clips, one per beat, the same woman who narrates the film. Not the visitor's laptop voice.
- The device voice survives only as a **fallback**, and the page says so out loud. Nothing is hidden.
- The character is a **guide, never the authority**: she explains, the AI proposes, the kernel decides, the human controls.
- Why it matters: a control layer is an abstract promise. A face and a voice make the promise **walkable** — and the same rule, said twice, is the cheapest trust there is.
- Walk it live: hackyeah-2026-ai-control-layer-production.up.railway.app/onboarding

What we measured — and what we do not claim.

MEASURED	VALUE	SOURCE
rows past the perimeter	0	checkpoint P2.3, fresh node
time to stop — live node	0.116 s	verify_live.py, 20/20
actions under a grant	100%	hash-chained receipts, recomputable from genesis
the whole suite	146 passed	pytest, on the shipped commit
this round, on camera	A-0002 · HTTP 200 · 19.5 ms	receipt 4db4d785, live node
embedding in a bank	1 endpoint · ~2 weeks	architecture: POST /mcp in front of the tool server

- **Not claimed:** a platform, an admin UI, a multi-tenant issuer. One node, one vector, proven.
- **Honest gap:** buyer willingness to pay is a hypothesis — no paid customer yet. The price is a hypothesis, not an invoice.